


```
# Import necessary libraries
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
```

```
# Load the dataset
df = pd.read_csv('/content/sonar data.csv')
```


```
df.head()
```




| | 0.0200 | 0.0371 | 0.0428 | 0.0207 | 0.0954 | 0.0986 | 0.1539 | 0.1601 | 0.3109 | 0.2111 | ... | 0.0027 | 0.0065 | 0.0159 | 0.0072 | 0.0167 | 0.0180 | 0. |
|---|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|-----|--------|--------|--------|--------|--------|--------|----|
| 0 | 0.0453 | 0.0523 | 0.0843 | 0.0689 | 0.1183 | 0.2583 | 0.2156 | 0.3481 | 0.3337 | 0.2872 | ... | 0.0084 | 0.0089 | 0.0048 | 0.0094 | 0.0191 | 0.0140 | 0. |
| 1 | 0.0262 | 0.0582 | 0.1099 | 0.1083 | 0.0974 | 0.2280 | 0.2431 | 0.3771 | 0.5598 | 0.6194 | ... | 0.0232 | 0.0166 | 0.0095 | 0.0180 | 0.0244 | 0.0316 | 0. |
| 2 | 0.0100 | 0.0171 | 0.0623 | 0.0205 | 0.0205 | 0.0368 | 0.1098 | 0.1276 | 0.0598 | 0.1264 | ... | 0.0121 | 0.0036 | 0.0150 | 0.0085 | 0.0073 | 0.0050 | 0. |
| 3 | 0.0762 | 0.0666 | 0.0481 | 0.0394 | 0.0590 | 0.0649 | 0.1209 | 0.2467 | 0.3564 | 0.4459 | ... | 0.0031 | 0.0054 | 0.0105 | 0.0110 | 0.0015 | 0.0072 | 0. |
| 4 | 0.0286 | 0.0453 | 0.0277 | 0.0174 | 0.0384 | 0.0990 | 0.1201 | 0.1833 | 0.2105 | 0.3039 | ... | 0.0045 | 0.0014 | 0.0038 | 0.0013 | 0.0089 | 0.0057 | 0. |

5 rows x 61 columns

```
print("Shape of the dataset:", data.shape)
```

 Shape of the dataset: (208, 61)

```
print(data.describe())
```



| | 0 | 1 | 2 | 3 | 4 | 5 | \ |
|-------|------------|------------|------------|------------|------------|------------|---|
| count | 208.000000 | 208.000000 | 208.000000 | 208.000000 | 208.000000 | 208.000000 | |
| mean | 0.029164 | 0.038437 | 0.043832 | 0.053892 | 0.075202 | 0.104570 | |
| std | 0.022991 | 0.032960 | 0.038428 | 0.046528 | 0.055552 | 0.059105 | |
| min | 0.001500 | 0.000600 | 0.001500 | 0.005800 | 0.006700 | 0.010200 | |
| 25% | 0.013350 | 0.016450 | 0.018950 | 0.024375 | 0.038050 | 0.067025 | |
| 50% | 0.022800 | 0.030800 | 0.034300 | 0.044050 | 0.062500 | 0.092150 | |
| 75% | 0.035550 | 0.047950 | 0.057950 | 0.064500 | 0.100275 | 0.134125 | |
| max | 0.137100 | 0.233900 | 0.305900 | 0.426400 | 0.401000 | 0.382300 | |

| | 6 | 7 | 8 | 9 | ... | 50 | \ |
|-------|------------|------------|------------|------------|-----|------------|---|
| count | 208.000000 | 208.000000 | 208.000000 | 208.000000 | ... | 208.000000 | |
| mean | 0.121747 | 0.134799 | 0.178003 | 0.208259 | ... | 0.016069 | |
| std | 0.061788 | 0.085152 | 0.118387 | 0.134416 | ... | 0.012008 | |
| min | 0.003300 | 0.005500 | 0.007500 | 0.011300 | ... | 0.000000 | |
| 25% | 0.008090 | 0.008425 | 0.0097025 | 0.011275 | ... | 0.008425 | |
| 50% | 0.106950 | 0.112100 | 0.152250 | 0.182400 | ... | 0.013900 | |
| 75% | 0.154000 | 0.169600 | 0.233425 | 0.268700 | ... | 0.020825 | |
| max | 0.372900 | 0.459000 | 0.682800 | 0.710600 | ... | 0.100400 | |

| | 51 | 52 | 53 | 54 | 55 | 56 | \ |
|-------|------------|------------|------------|------------|------------|------------|---|
| count | 208.000000 | 208.000000 | 208.000000 | 208.000000 | 208.000000 | 208.000000 | |
| mean | 0.013420 | 0.010709 | 0.010941 | 0.009290 | 0.008222 | 0.007820 | |
| std | 0.009634 | 0.007060 | 0.007301 | 0.007088 | 0.005736 | 0.005785 | |
| min | 0.000800 | 0.000500 | 0.001000 | 0.000600 | 0.000400 | 0.000300 | |
| 25% | 0.007275 | 0.005075 | 0.005375 | 0.004150 | 0.004400 | 0.003700 | |
| 50% | 0.011400 | 0.009550 | 0.009300 | 0.007500 | 0.006850 | 0.005950 | |
| 75% | 0.016725 | 0.014900 | 0.014500 | 0.012100 | 0.010575 | 0.010425 | |
| max | 0.070900 | 0.039000 | 0.035200 | 0.044700 | 0.039400 | 0.035500 | |

| | 57 | 58 | 59 |
|-------|------------|------------|------------|
| count | 208.000000 | 208.000000 | 208.000000 |
| mean | 0.007949 | 0.007941 | 0.006507 |
| std | 0.006470 | 0.006181 | 0.005031 |
| min | 0.000300 | 0.000100 | 0.000600 |
| 25% | 0.003600 | 0.003675 | 0.003100 |
| 50% | 0.005800 | 0.006400 | 0.005300 |
| 75% | 0.010350 | 0.010325 | 0.008525 |
| max | 0.044000 | 0.036400 | 0.043900 |

```
[8 rows x 60 columns]
```

2 Data Preparation

```
X = df.iloc[:, :-1] # Features
y = df.iloc[:, -1] # Target
```

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y)
```


```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

3 Model Development

```
from sklearn.neighbors import KNeighborsClassifier
neighbors = np.arange(1, 14)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))
```

```
for i, k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    train_accuracy[i] = knn.score(X_train, y_train)
    test_accuracy[i] = knn.score(X_test, y_test)
```

```
plt.plot(neighbors, test_accuracy, label='Testing Accuracy')
plt.plot(neighbors, train_accuracy, label='Training Accuracy')
plt.xlabel('Number of Neighbors')
plt.ylabel('Accuracy')
plt.title('k-NN Varying Number of Neighbors')
plt.legend()
plt.show()
```



```
knn = KNeighborsClassifier(n_neighbors=2)
knn.fit(X_train, y_train)
y_pred_knn = knn.predict(X_test)
```

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
y_pred_logistic = model.predict(X_test)
```

```
# Initialize PCA and reduce the number of components
pca = PCA(n_components=10) # Adjust the number of components as needed
```

```

X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Train Logistic Regression model on PCA-transformed data
model_pca = LogisticRegression()
model_pca.fit(X_train_pca, y_train)

# Make predictions
y_pred_pca = model_pca.predict(X_test_pca)

# Train SVM model on original features
svm = SVC(kernel='linear') # You can use different kernels like 'rbf' as well
svm.fit(X_train, y_train)

# Make predictions
y_pred_svm = svm.predict(X_test)

```

4 Model Evaluation

```

from sklearn.metrics import accuracy_score, confusion_matrix
print("kNN Accuracy:", accuracy_score(y_test, y_pred_knn))

```

```

➦ kNN Accuracy: 0.7619047619047619

```

```

print("kNN Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_knn))

```

```

➦ kNN Confusion Matrix:
[[31  5]
 [10 17]]

```

```

print("Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_logistic))

```

```

➦ Logistic Regression Accuracy: 0.746031746031746

```

```

print("Logistic Regression Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_logistic))

```

```

➦ Logistic Regression Confusion Matrix:
[[24 12]
 [ 4 23]]

```

```

# Evaluate the PCA-based model
print("PCA + Logistic Regression Accuracy:", accuracy_score(y_test, y_pred_pca))
print("PCA + Logistic Regression Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_pca))

```

```

➦ PCA + Logistic Regression Accuracy: 0.7619047619047619
PCA + Logistic Regression Confusion Matrix:
[[24 12]
 [ 3 24]]

```

```

# Evaluate the SVM model
print("SVM Accuracy:", accuracy_score(y_test, y_pred_svm))
print("SVM Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_svm))

```

```

➦ SVM Accuracy: 0.7619047619047619
SVM Confusion Matrix:
[[26 10]
 [ 5 22]]

```

Generated code may be subject to a license | CallmeAk/Data-Analyst-Projects | AnmolGulati6/Rock-vs-Mines-Classification-using-Logistic-Regression

```
input_data = (0.0100,0.0171,0.0623,0.0205,0.0205,0.0368,0.1098,0.1276,0.0598,0.1264,0.0881,0.1992,0.0184,0.2261,0.1729,0.2131,0.0693,0.2281,0
```

```

#changing the input_data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)

```

```

#reshape the np array as we are predicting for one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)


```

```

prediction = model.predict(input_data_reshaped)
print(prediction)

```

```
if (prediction[0]=='R'):  
    print('The object is a Rock')  
else:  
    print('The object is a mine')
```

```
↵ [1]  
The object is a mine  
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names, but Logistic  
warnings.warn(  

```